

Application Security (apsi)

Lecture at FHNW

Lecture 2, 2021

Arno Wagner, Michael Schläpfer

<arno@wagner.name>, <{michael.schlaepfer, rolf.wagner}@fort-it.ch>

Agenda

More on attack techniques and vulnerabilities

- ▶ Buffer Overflow II
- ▶ Error handling: The good, the bad and the ugly
- ▶ Algorithmic complexity attacks

What do Attackers Want?

- ▶ Access to data on the target
- ▶ Manipulation of data on the target
- ▶ Sabotage of a service on a system (political goals, extortion, vandalism)
- ▶ Identity (IP address, DNS name): SPAM, DDoS, jump-off for further attacks
→ The last item is why "hack-back" is a really, really bad idea...
- ▶ Resources (RAM, CPU....)
Newer Trend: Crypto-currency mining

Typical means to get these: Compromise user or system account

Other approaches: Compromise Browser, VM, Sandbox (container, e.g. "Docker") as execution environment
→ Isolation of code may not help! (As long as there is network access...)

Buffer Overflow II

Note: The whole thing is an arms-race. No end is in sight. The defenders are not winning. Robust code (correct and/or redundant) is the only reliable fix.

Now: Buffer overflow used to overwrite variables:

- ▶ Also works on heap, in shared-memory, on-disk data-structures, etc.
→ Everywhere where values of variables are stored!

Extensions:

- ▶ Code execution via buffer overflow
Note: May or may not require code injection
"Return-oriented programming" does not
- ▶ Finding buffer-overflow bugs as an attacker via "Fuzzing"

Buffer Overflow II

```
#include <string.h>

#include <stdio.h>

void test (int argc, char *argv[]) {
    char s2[4] = "yes";    // set s2 to "yes"
    char s1[4] = "abc";    // set s1 to "abc"
    strcpy(s1, argv[1]);    // copy argv[1] into s1
    puts(s2);              // print s2
}

int main (int argc, char *argv[]) {
    test(argc, argv);
}
```

Demo with GDB: Make input string much larger. What happens?

Buffer Overflow II

Stackframe also contains return address

Note: On x86 this is an absolute address.

- ▶ Can be overwritten
 - This usually causes a segfault on function return as most addresses are invalid or non-executable memory.
- ▶ Can be overwritten in a targeted fashion
 - ▶ Jump to attack-code in same buffer
 - ▶ Jump to attack-code in a different buffer
 - ▶ Jump to some library-call ("Return-oriented Programming")

Buffer Overflow II

Example OS-Level countermeasure: Memory-layout randomization

Idea:

- ▶ Place stack and heap at random places so attacker does not know where the inserted attack code is and hence where to jump

Countermeasure:

- 1) "NOP-slide"/"Jump catcher": Large area of valid entry-points in attack code
- 2) Jump to random location

Note: Catch-area can be made very large

→ Randomization usually only makes the attack somewhat more difficult

Buffer Overflow II

OS-Level countermeasure: NX-Bit

Idea:

- ▶ Prevent code execution on the stack, the heap and other non-code areas

Countermeasure: "return-oriented programming"

- ▶ Put in jumps to a sequence of legitimate library routines and code fragments
- ▶ When one returns, the next one is called
- ▶ This technique is generally Turing-complete

→ NX often only helps to make attack more difficult

More detail:

https://www.blackhat.com/presentations/bh-usa-08/Shacham/BH_US_08_Shacham_Return_Oriented_Programming.pdf

Attack Technique: Fuzzing

General technique to find buffer-overflows (and other errors)

→ Also useful to test some types of code

Idea:

- ▶ Throw random values (unstructured fuzzing) or randomized structured values (structured fuzzing, e.g. valid XML or JSON) at an interface
- ▶ Check whether the software crashed or behaves unexpectedly

→ This identifies potential points of attack

Characteristics:

- ▶ Easy to automatize
- ▶ Finds buffer overflows and bad error handling fast
- ▶ Limited power for complex interactions (e.g. login)

Error Handling

Error handling is a vulnerability if the error behavior is problematic

- ▶ DoS by resource exhaustion (RAM, CPU, I/O, etc.)
- ▶ Data leakage (user names, account numbers, etc. → Lecture 1)
- ▶ Leakage of software and OS versions
- ▶ Leakage of internals
- ▶ Hints that a buffer-overflow or other problem may exist
- ▶ Hints to shoddy coding practices
- ▶ ...

Case study – Handling too long input

Assume your server gets a query which is too long. What can you do?

- ▶ Verbose error
- ▶ Generic error
- ▶ Silent failure
- ▶ Cut it down (dumb or smart)
- ▶ Try to accomodate anyways
 - ▶ For example: only if space, slow it down, ask for (more) money, ...
- ▶ Transform it: Compress, change encoding (picture), remove markup,...
- ▶ Ask user to adjust
- ▶ Ask user what to do

Remember that the limitations and strategy used will be pretty obvious!

All have serious problems... You still need to select one.

Example Attack: Compression Bomb

A server accepts compressed input and scans it (e.g. Email anti-virus)

- ▶ Such data can expand to basically arbitrary length!
- ▶ Server cannot drop it (may be important), cannot ask for better input, etc.
- ▶ Cannot decompress it before processing, too large

What do do?

Possible solution: Stream-processing

- ▶ Decompress incrementally, scan, recompress
- ▶ May still take a long time, but no full disks, no exhausted memory, etc.
- ▶ It may be acceptable to fail silently after a lot of data (judgment call...)
But: Email size limits are a problem for large files, e.g. Word documents

Error Handling Practices (1)

Bad practices:

- ▶ Fail silently (unless you are really sure it is an attack)
- ▶ Give out too little or too much information
- ▶ Call debug code in production (oops...)
- ▶ Consume a lot of resources
- ▶ Log more error-data than you can accomodate
→ Attacker may first flood logs and then experiment unobserved
- ▶ Have side-effects on other activities
- ▶ Give misleading information (does not really deter attackers)
- ▶ Make it obvious that some errors are not handled (for example by crashing)
- ▶ ...
- ▶ Anything overly complex and violating KISS

Error Handling Practices (2)

Good practices (exceptions may apply, document if you deviate):

- ▶ Handle all errors
- ▶ Be fault-tolerant whenever it does not cost you much (but prevent DoS)
- ▶ Be user-friendly unless that leaks data or can cause DoS
- ▶ Do not give details that do not refer to the input
- ▶ Do not echo the input (prevent reflector attacks)
- ▶ Test error code under high error-load
- ▶ Mimimize extra resources used in error handling
- ▶ Expect error handling to be abused
- ▶ ...
- ▶ And always: KISS!

Algorithmic Complexity Attacks

Attack type: DoS (i.e. sabotage), may also be used as supporting attack

- ▶ Exploits bad worst-case complexity of algorithms that usually perform well
 - ▶ Many algorithms commonly used today have this property!
- ▶ Aims at massive slowdown or complete unavailability
 - ▶ Note: Unavailability can also be a result of safety mechanisms
Example: The banking-industry often aborts Mainframe-queries after 500ms
- ▶ Effective protection usually exists, but awareness is critical
... as so often in secure software engineering and secure systems...

Algorithmic Complexity: Sorting

Sorting (the "classical" example):

Quicksort is still commonly used

- ▶ It has $O(n \log n)$ space (memory) and time complexity on average
- ▶ It takes $O(n^2)$ space (!) and time in the worst case (space is often on the stack)
- ▶ Worst-case is easy to hit: Just feed it pre-sorted data

Fix: Do not use the obsolete Quick-Sort. No, really not!

Use merge-sort or bottom-up heapsort.

=> Minimally slower, but worst-case optimal in time and space.

Algorithmic Complexity: Hashing

Hash-tables are widely used for lookup

- ▶ Average case complexity $O(1)$
- ▶ Worst-case complexity: $O(n)$ (time)

Recap: Idea of a hash-table (Note: h is not a cryptographic hash):

- ▶ Store value at a position determined by a hash-function $h(\text{value})$ in an array
→ Hash-function should distribute elements well
- ▶ Collision when $h(\text{value1}) == h(\text{value2})$: Chain colliding elements in some form
- ▶ Attack: Select elements that will cause a lot of collisions

Fix: Add a secret key-quality random element r (≥ 64 bit) to each table instance.
Use $h(r, \text{value})$ as hash-function.
Also use a modern hash-function like Spooky-Hash or Sip-Hash!

Algorithmic Complexity: Regular Expressions

Regular expression can have exponential execution time (input length)

Attack strategies:

- ▶ Use such a string on a bad regular expression (common)
- ▶ Supply a bad regular expression (and string to trigger it) to the target (rare)

Simple Example:

- ▶ $(a^+)^+$ with greedy matching, "aaa...aaaab" as string.

Fix:

- ▶ Prevent use of such regular expressions
- ▶ Limit runtime and go into error handling if exceeded

Note: Vulnerability depends on matching-engine. NFA-Engines and lazy matching are not vulnerable. These have limitations though, for example no capturing groups in NFA engines.