

Application Security (apsi)

Lecture at FHNW

Lecture 3, 2021

Arno Wagner, Michael Schläpfer, Rolf Wagner

<arno@wagner.name>, <{michael.schlaepfer,rolf-wagner}@fort-it.ch>

Agenda

- ▶ Test-before-use, Input Validation
- ▶ Input Normalization
- ▶ Privilege models in Linux (Unix): file access, running processes
- ▶ Chroot(2)
- ▶ Mandatory Access Control (MAC): SELinux, AppArmor
- ▶ Architecture-Pattern: Privilege Separation

Secure Coding Principle: Test-before-use

This is a fundamental secure coding principle

- ▶ Always test all assumptions about data before using it in any way
→ This makes assumptions explicit and documents them
- ▶ Code only has to work for the tested properties
→ This simplifies the code and helps with KISS
- ▶ Examples
 - ▶ Length of strings and buffers
 - ▶ Allowed characters in strings
 - ▶ Numerical ranges
 - ▶ State of sockets or files
 - ▶ ...

Input Validation

Purpose: Establish properties of input data first

→ Instance of the "test-before-use" principle

- ▶ Later processing stages can then depend on the tested properties
→ This makes coding and error handling much easier.

What needs to be validated?

- ▶ *Every* property of input data the program depends on!
→ You may only assume input data properties you have validated first!
(Insecure software is almost always caused by invalid assumptions...)

▶ Examples:

- ▶ Input length (min, max) → Buffer overflow. "Too short" can also be a problem.
- ▶ Input character set, escaping or not, special codes (for example \0)
- ▶ Other formats, e.g. "valid XML", "valid JSON", "valid HTML", etc.
- ▶ Configuration files

Input Normalization

This is a preprocessing step, used to decrease variability of input data

The goal is simplification of further processing (again "KISS")

- ▶ Reduce a larger set of input formats to a smaller one or to a single format
- ▶ Simplifies following format-validations and processing
- ▶ May have input constraints as well, input validation before still needed

Examples:

- ▶ Decompress data or images for multiple compressors and (optionally) recompress with a specific one
- ▶ Recode all text-input to UTF-8
- ▶ Virtually print a document, use OCR on it to recover the text
→ This is currently the only really secure way to sanitize documents.

Unix Permission Model (simplified)

- 1) Everything is a file (includes sockets, directories, interfaces, etc.)
- 2) Every file has 3 permission sets: user, group, other
- 3) Every file can be (r)readable, (w)ritable and e(x)ecutable in each group
Note: For directories "x" is required for chdir-ing into them.
"x" becomes "s" on executable files with suid/sgid bit set
- 4) A running process has a user and a group which determines its rights
 - ▶ This is simple, but often enough. If not → Mandatory Access Control (MAC)
 - ▶ "root" is special and can change all permissions
 - ▶ A process running as root can "drop privileges" to the permissions of a regular user by calling `setgid()` and `setuid()`. (Done right, it cannot get back.)
 - ▶ The suid/sgid-bit executes that file as the user/group it belongs to.
→ This allows ordinary users to execute programs as root
Example: `/bin/su: -rwsr-xr-x 1 root root 36816 May 25 2012 /bin/su`

Chroot

Sets the apparent root-directory ("/") for a process

- ▶ Isolation technique
- ▶ If done right, process cannot access files outside
- ▶ Usually used before a "privilege-drop"
- ▶ Can be used to "trap" a running process in a subdirectory

Limitation:

- ▶ Root can break out of a chroot situation

Note: There is both a C call `chroot(2)` and a shell-call `chroot(8)`

Reference:

- <http://www.unixwiz.net/techtips/chroot-practices.html>

Mandatory Access Control (MAC)

Fine-grained access control enforced by the OS kernel

- ▶ Allows very specific permissions
- ▶ Allows restricting "root" (usually needs reboot to change)
- ▶ Knows OS objects are more than just files (sockets, directories, etc.)
- ▶ Has a concept of groups, allows control over entering and leaving groups
- ▶ Cannot be switched off at run-time (unless specifically allowed)
- ▶ Allows a "fully locked down" system, even against root.

But:

- ▶ Difficult to configure
- ▶ Can cause hard to debug run-time problems
- ▶ May hinder normal system administration

MAC Example: SELinux

- ▶ Initial release: 1998, created by the NSA (this should not be a problem)
- ▶ Does labelling for everything in the file-system
→ can take a while on first start
Needs filesystem that supports security-labels
- ▶ 3 Modes: enforcing, permissive, disabled
- ▶ Can restrict what root can do, can prevent root login
- ▶ Complex to configure and maintain
- ▶ Very fine-grained control

Reference: https://en.wikipedia.org/wiki/Security-Enhanced_Linux

MAC Example: AppArmor

- ▶ Initial release: 1998
- ▶ Works with file-paths (i.e. not files directly)
- ▶ Simple (relatively to SELinux) to configure
- ▶ Does not need a file-system with security labels
- ▶ Can be switched-off on a per-path level
- ▶ Less fine control than SELinux
- ▶ Less structured than SELinux

Reference: <https://en.wikipedia.org/wiki/AppArmor>

Secure Coding Principle: Defense in Depth

Architecture and coding principle

Idea: Have several security mechanisms so that one is enough to still be secure if the others have been circumvented by an attacker.

- ▶ Example: Input validation and least privilege
- ▶ Example: Password and smartcard

Rationale:

- ▶ Gives resilience if one mechanism fails (and mechanisms will fail...)
- ▶ Gives resilience against bugs in the implementation of mechanisms
- ▶ Can help in attack-detection
- ▶ Allows individually less-complicated protection mechanisms (KISS)
→ Complex protection mechanisms can be a problem! Example: IPSec
- ▶ Makes attacks more expensive → Discourages attackers

Privilege Separation

- ▶ Consequence of the Principles of "Least privilege" and "Defense in depth"
Idea: Run code only with the minimal privileges needed

But:

- ▶ Code is complex and different parts need different privileges

Hence:

- ▶ Separate code into modules, run each one with minimal privileges needed
- ▶ In particular separate things like input normalization, network and system access, user interaction, etc.
 - Modules later in the data-flow are not directly exposed
 - If data is incorrect later, refuse processing

This technique has allowed some very exposed server software to survive well in a very hostile Internet

Overall Idea

Two main approaches (will be combined):

1) Operations that require high privilege (root) are done before any client interaction, then drop privileges before procesing data

- ▶ Example: Allocating a network socket (ports below 1024 require root-rights)
- ▶ Example: Accept a connection but then hand it to a non-privileged child
- ▶ Example: Allocating non-swappable memory

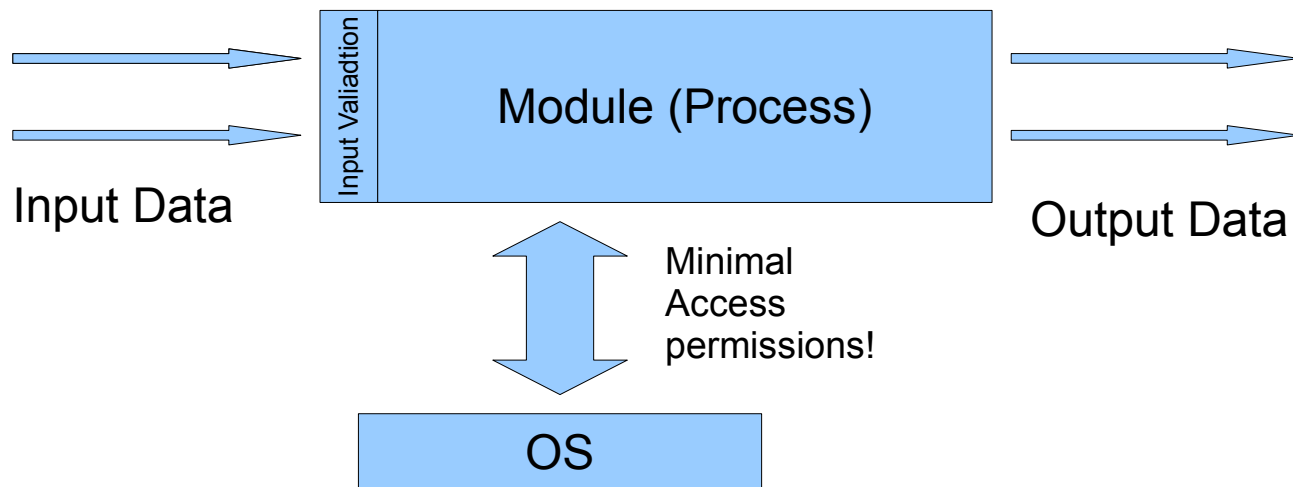
2) Divide processing into modules along the data-path

- ▶ Input validation first and again in each module
- ▶ Normalization next
- ▶ Actual processing next

Communication between elements: Files and usual Unix IPC like local sockets, pipelines, signals, ...

Form of the Modules

- ▶ A module is encapsulated into one or more processes (or machine, container, etc. for larger security needs)
 - Run each module with the minimal access needed
 - If higher privileges needed for some steps, do them before processing input data and then drop privileges
- ▶ Do full input validation in each module!
- ▶ If module seems to require 2 sets of privileges for data processing, split it!



Principles for Privilege Separation

- ▶ Always do a full (!) input validation on any internal communication channel
- ▶ Keep internal communication as simple as possible (see "KISS")
-> This makes input validation simpler and reduces vulnerabilities in it
- ▶ If input validation fails at an internal channel, if possible log the issue and optionally find the sending component to be compromised. (Application-internal intrusion detection.) Consider preparing countermeasures.
- ▶ Always make sure compromising another module from a compromised one is as hard as possible
- ▶ Do not expose the internal communication
- ▶ Treat configuration files, data files, etc. as communication channels
- ▶ Select the right level of separation: High security requirements need separation in small components, low security requirements can use larger, less restricted components.

Suitable and Unsuitable Communication Mechanisms

Suitable for communication between privilege-separated modules:

- ▶ Anonymous and named pipes, local sockets ("Unix Domain Sockets")
- ▶ Signals
- ▶ Files in the Filesystem (with proper restrictions by permissions)

Note: If forking children: Before you do anything else:

1. Delete all file-descriptors from the parent that the child does not need
2. Erase all memory that contains any secrets of the parent

Unsuitable mechanisms:

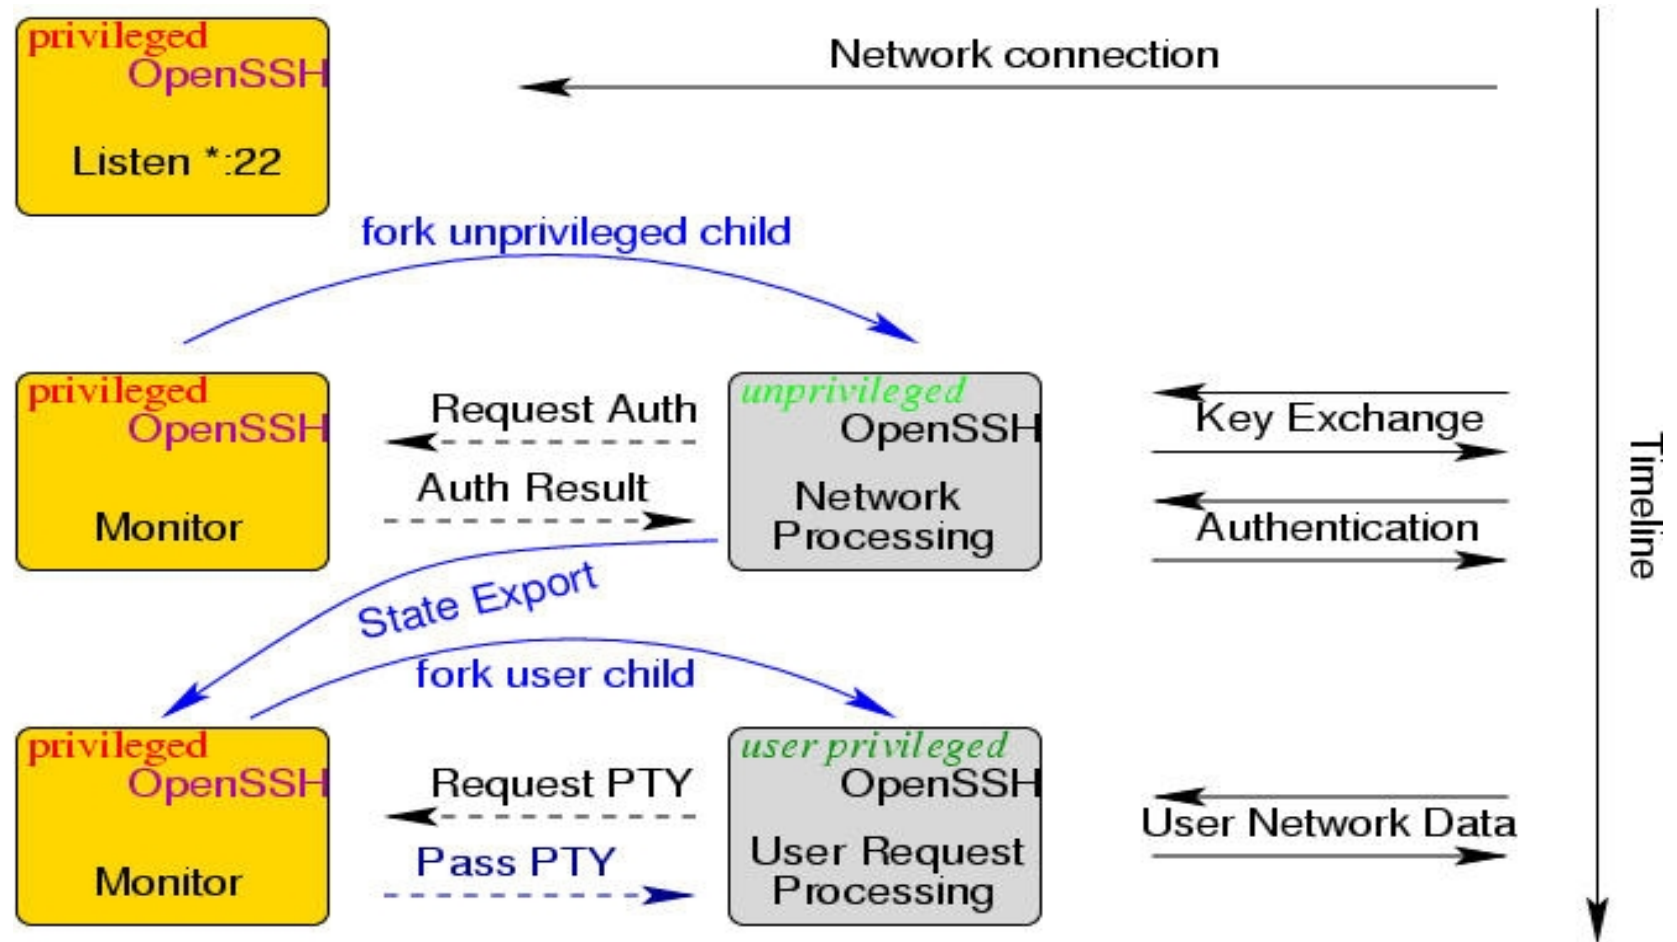
- ▶ RPC (Remote Procedure Calls): Have a very bad security track record
- ▶ Shared memory: A buffer overflow can propagate...
- ▶ Network sockets: You expose internal communication to the network!
If using 127.0.0.1, you still expose internal communication to other processes!
- ▶ Sending serialized objects: You may end up sending compromised objects!

Example: OpenSSH

This is the OpenBSD Secure Shell Server

- ▶ Initial release 1999
- ▶ The standard way for secure shell login to a UNIX machine
- ▶ Obviously an attractive target
- ▶ Must be Internet-reachable in most cases
- ▶ Must protect private keys used for authentication
- ▶ Details: <http://www.citi.umich.edu/u/provos/ssh/privsep.html>

OpenSSH: Privilege Separation Approach



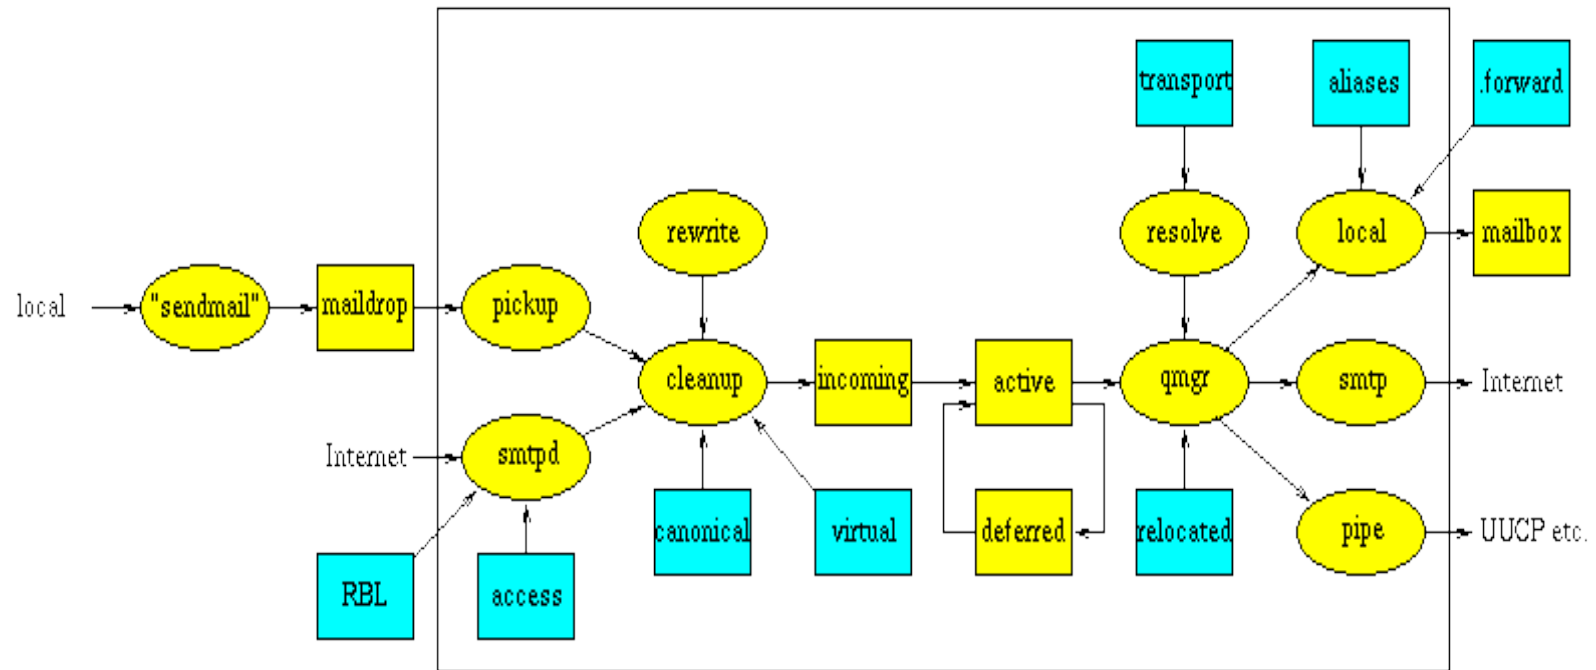
Example: Postfix MTA (Email Server)

The Postfix secure Email Server

- ▶ Initial release 1998
- ▶ Large, complex Email handling server (smtpd)
- ▶ Must be Internet-Reachable without any authentication or limit as to source
Possible exception: SPAM-sources
- ▶ Must open connections to arbitrary other MTAs
- ▶ "In April 2016 in a study performed by E-Soft, Inc., approximately 32% of the publicly reachable mail-servers on the Internet ran Postfix."
- ▶ Email is often a critical communication channel

Reference: <http://www.postfix.org/documentation.html>

Postfix System View



Postfix: Example Data-Flow

Incomming email → local user

smtpd(8) -> cleanup(8) -> incoming -> qmgr(8) -> local(8)

- ▶ **smtpd**: Accepts email from the Internet.
Filehandles/sockets and pipeline for giving mail to cleanup is inherited.
Process has no permissions at all.
- ▶ **cleanup**: Validates and normalizes email-format.
Is chroot'ed on the email-queue-directory, to which it can only write, not read.
Is allowed to send signals to **qmgr**.
- ▶ **qmgr**: Decides where to send (local, send over the net, pipe to command).
Only reads email header, not content or attachments.
Is chroot-ed on several mail-queue directories.
Allowed to signal the delivery demons **local**, **smtp**, **pipe**
- ▶ **local**: Does local delivery. Runs as root.
Gets target-user file-name in queue-directory via pipe.
Moves file to user's mailnoX, but never reads it.